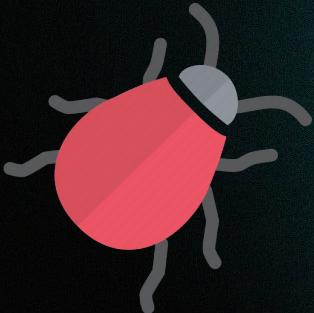


Introduction To Fuzzing



```
american fuzzy lop ++2.60d (libpng) [explore] {0}
process timing
  run time : 0 days, 0 hrs, 4 min, 43 sec
  last new path : 0 days, 0 hrs, 0 min, 0 sec
  last uniq crash : none seen yet
  last uniq hang : none seen yet
cycle progress
  now processing : 127.0 (60.2%)
  paths timed out : 0 (0.00%)
stage progress
  now trying : arith 16/8
  stage execs : 102/6963 (1.46%)
  total execs : 1.39M
  exec speed : 5147/sec
fuzzing strategy yields
  bit flips : 19/51.9k, 8/51.9k, 1/51.8k
  byte flips : 0/6488, 0/5262, 0/5210
  arithmetics : 7/295k, 1/114k, 0/52.1k
  known ints : 1/27.0k, 1/122k, 0/204k
  dictionary : 0/0, 0/0, 3/257k
  havoc/rad : 23/114k, 0/0, 0/0
  custom/rq : 0/0, 0/0, 45/9658, 12/15.2k
  trim : 36.12%/3042, 18.23%
overall results
  cycles done : 0
  total paths : 211
  uniq crashes : 0
  uniq hangs : 0
map coverage
  map density : 0.24% / 1.67%
  count coverage : 1.81 bits/tuple
findings in depth
  favored paths : 108 (51.18%)
  new edges on : 126 (59.72%)
  total crashes : 0 (0 unique)
  total tmouts : 0 (0 unique)
path geometry
  levels : 3
  pending : 166
  pend fav : 61
  own finds : 121
  imported : 86
  stability : 100.00%
[cpu000: 31%]
```



About Me

- What I do:
 - Ph.D. student at the University of Utah
 - Performing research in cybersecurity
 - Uncovered over 50 bugs in software
 - Talk to me about the program!
- Other things about me:
 - Love being outside
 - Have a weiner dog

Principal Investigator

Dr. Stefan Nagy
Assistant Professor

Graduate Student Researchers

Zao Yang
PhD Student

Yaseen Asafat
PhD Student

Gabe Sherman
PhD Student

Undergraduate Student Researchers

Josh Dean

Dillon Otto

Raj Reddy

Tanner Rowlett

A person snowboarding down a snowy mountain slope, giving a thumbs up.

Check out our lab: <https://futures.cs.utah.edu/>

Software Is Buggy!

This can have disastrous effects ...

Ballista Botnet Exploits Unpatched TP-Link Vulnerability, Infects Over 6,000 Devices

Mar 11, 2025 · Ravie Lakshmanan

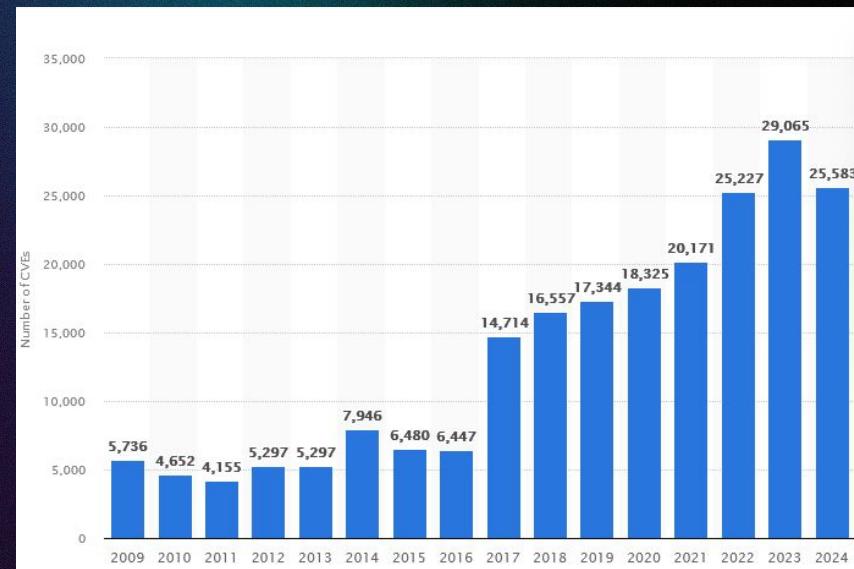
The ‘most serious’ security breach ever is unfolding right now. Here’s what you need to know.

Dec. 20, 2021 at 10:00 pm | Updated Dec. 20, 2021 at 10:25 pm



Heartbleed Bug: Tech firms urge password reset

① 9 April 2014



<https://www.statista.com/statistics/500755/worldwide-common-vulnerabilities-and-exposures/>

Software Uses Inputs

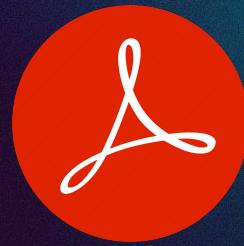
Audio



Video



Photos



And more...



When these inputs are not properly handled, bugs occur!

How do we find bugs?

Manual Analysis

- Codebases are huge! (Millions of lines of code)
- Massive amounts of effort
- Still might miss something

Static Analysis

- Analyze program without running it
- Quickly becomes computationally infeasible
- False positives and false negatives

Dynamic Analysis

- Analyze program by running it!
- What you see is what you get
- Minimal overhead



Fuzzing

A dynamic analysis technique that aims to execute a program as many times as possible with random inputs.

- Widely used in today's software ecosystem
- Google's OSS-Fuzz project has found over 10,000 vulnerabilities and 36,000 bugs through fuzzing

Fuzzing Reveals Over 30 Web Browser Engine Flaws

Fuzzing tests conducted on the most popular web browser engines by Google Project Zero revealed the existence of more than 30 vulnerabilities, more than half of which in Apple's Safari.

New fuzzing tool finds 26 USB bugs in Linux, Windows, MacOS, and FreeBSD

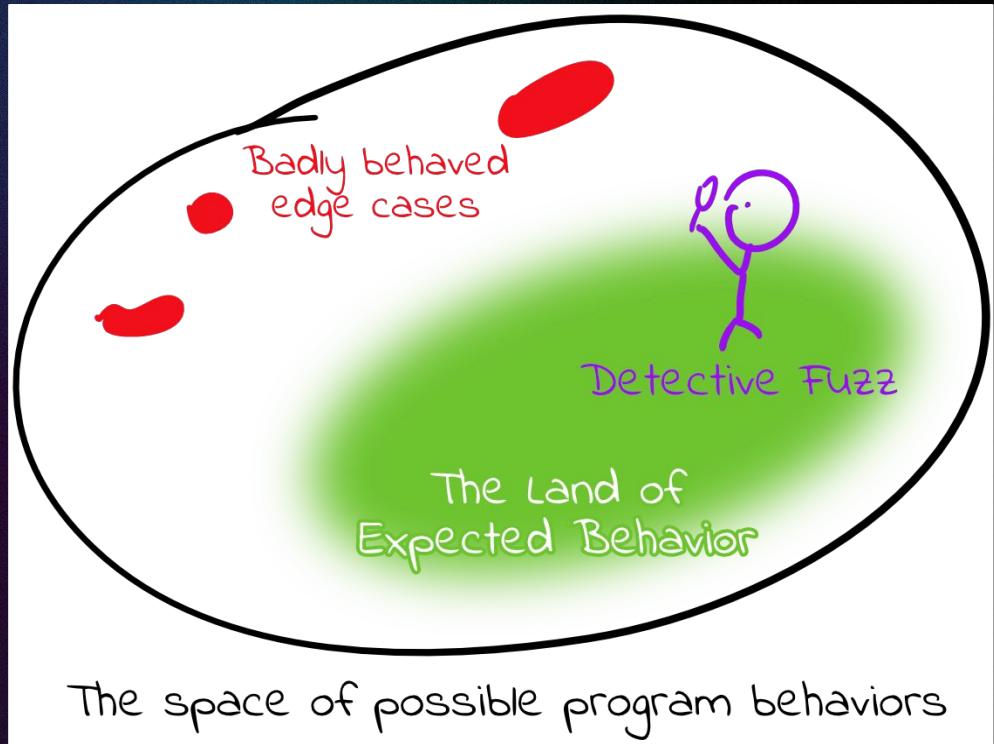
Eighteen of the 26 bugs impact Linux. Eleven have been patched already.

Understanding Shellshock [CVE-2014-6271]: A Critical Bash Vulnerability

Fuzzing

The General Idea:

- Quickly generate a ton of random inputs
- Execute the program on each input. Does the program:
 - Crash?
 - Hang?
 - Do Nothing?



What can you fuzz?

Everything

What can you fuzz?



JANUS: Detecting Rendering Bugs in Web Browsers via Visual Delta Consistency

Chijin Zhou[†], Quan Zhang[†], Bingzhou Qian[†], and Yu Jiang^{‡✉}

[†]BNRist, Tsinghua University, Beijing, China

[‡]National University of Defense Technology, Changsha, China

FUZZILLI: Fuzzing for JavaScript JIT Compiler Vulnerabilities

Samuel Groß
Google
saleo@google.com

Simon Koch
TU Braunschweig
simon.koch@tu-braunschweig.de

Lukas Bernhard
Ruhr University Bochum
lukas.bernhard@rub.de

Thorsten Holz
CISPA Helmholtz Center for Information Security
holz@cispa.de

Martin Johns
TU Braunschweig
m.johns@tu-braunschweig.de

What can you fuzz?



WHITEFOX: White-Box Compiler Fuzzing Empowered by Large Language Models

CHENYUAN YANG, University of Illinois at Urbana-Champaign, USA

YINLIN DENG, University of Illinois at Urbana-Champaign, USA

RUNYU LU, Huazhong University of Science and Technology, China

JIAYI YAO, Chinese University of Hong Kong, China

JIWEI LIU, University of Illinois at Urbana-Champaign, USA

REYHANEH JABBARVAND, University of Illinois at Urbana-Champaign, USA

LINGMING ZHANG, University of Illinois at Urbana-Champaign, USA

Finding and Understanding Bugs in C Compilers

Xuejun Yang Yang Chen Eric Eide John Regehr

University of Utah, School of Computing
{jxyang, chenyang, eeide, regehr}@cs.utah.edu

What can you fuzz?



Fuzzware: Using Precise MMIO Modeling for Effective Firmware Fuzzing

Tobias Scharnowski¹, Nils Bars¹, Moritz Schloegel¹, Eric Gustafson², Marius Muench³, Giovanni Vigna^{2,4}, Christopher Kruegel², Thorsten Holz¹ and Ali Abbasi¹

¹Ruhr-Universität Bochum, ²UC Santa Barbara, ³Vrije Universiteit Amsterdam, ⁴VMware

SNIPUZZ: Black-box Fuzzing of IoT Firmware via Message Snippet Inference

Xiaotao Feng
Swinburne University of Technology
Australia

Minhui Xue
The University of Adelaide
Australia

Sheng Wen*
Swinburne University of Technology
Australia

Ruoxi Sun
The University of Adelaide
Australia

Surya Nepal
CSIRO Data61
Australia

Yang Xiang
Swinburne University of Technology
Australia

Xiaogang Zhu*
Swinburne University of Technology
Australia

Dongxi Liu
CSIRO Data61
Australia

What can you fuzz?



OpenSSL



AFLplusplus

- *Coverage guided, grey-box fuzzer*
- Most widely used fuzzing engine out there (5.5k github stars)
- High portability and adaptability
- Learn how to use this when learning fuzzing!
- Some useful links:
 - <https://github.com/AFLplusplus/AFLplusplus>
 - <https://aflplus.plus/>



Coverage-guided fuzzing

- Keep and further mutate inputs that gain new coverage in the program under test
- No new coverage means we didn't trigger new functionality!

Fuzzers generate tons of inputs. How we determine if an input is interesting or meaningful?

```
4 177x function fib(n) {
5 177x   if (n === 0) {
6 34x     return 0
7 177x   } else if (n === 1) {
8 55x     return 1
9 143x   } else if (n > 1) {
10 88x     return fib(n - 1) + fib(n - 2)
11 } else {
12   thrower()
13 }
14 177x
15 1x console.log('fib(10):', fib(10))
```

A Simple Example



bgatradf

Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

New coverage! –
Mutate this input

A Simple Example



No new coverage

Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

A Simple Example



aertstbwrr

Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

New coverage! –
Mutate this input

A Simple Example



abwrertst

Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

New coverage! –
Mutate this input

A Simple Example



acwrertst

Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

Crash!

The Overall Process

Input Files



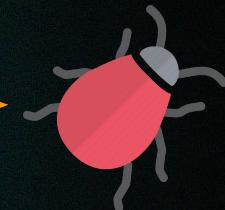
Execute program
and collect
feedback

New coverage



No new coverage

Crash!



More On Coverage

How do programs report their coverage?

- *Instrument* the program to report when regions of code are reached.
- Open Source?
 - Instrument at compilation time – afl-clang-fast and afl-clang-fast++ are good in most cases
- Closed Source?
 - Instrument on the fly using afl-qemu

Uninstrumented Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

Instrumented Program

```
if(input[0] == 'a'){
    if(input[1] == 'b'){
        return 2;
    }
    if(input[1] == 'c'){
        Crash!
    }
}
```

Hey AFL, I'm being executed!

Get a headstart

- Fuzzers generate inputs and mutate them. Give it a headstart by providing *seeds*.
- Fuzzing seeds are *examples* of inputs the program may expect. This lets the fuzzer operate on and mutate these rather than starting from scratch.
- Seeds are *program dependent*.
- Some examples
 - Adobe acrobat: Different PDF files
 - Spotify: Different audio files
 - Imessage: .db files



fuzzing-seeds

A centralized collection of various popular fuzzing seed corpora.

- `seeds_af1/` -- <https://github.com/AFLplusplus/AFLplusplus/tree/master/testcases>
- `seeds_ankou/` -- <https://github.com/SoftSec-KAIST/Ankou-Benchmark>
- `seeds_dvyukov` -- <https://github.com/dvyukov/go-fuzz-corpus>
- `seeds_mopt/` -- <https://github.com/puppet-meteor/MOpt-AFL/tree/master/seed%20sets>
- `seeds_strongcourage/` -- <https://github.com/strongcourage/fuzzing-corpus>
- `dicts_google/` -- <https://github.com/google/fuzzing/tree/master/dictionaries>

<https://github.com/FuturesLab/fuzzing-seeds/>

Library Fuzzing

Libraries are used *everywhere* in software. Bugs and vulnerabilities that affect them not only affect the library itself, but software that depends on it!

For example...

XZ-Utils Backdoor: Adversary placed a backdoor in the liblzma library. Granted them remote access to linux systems.

**Linux could have been
brought down by backdoor
found in widely used utility**

The malicious code modifies functions within a data compression library that is a foundational part of several Linux distributions

OpenSSL Heartbleed Vulnerability: Improper input validation allowed attackers to leak sensitive information (secret keys, passwords, etc). 500,000 servers were immediately vulnerable.



How to Fuzz a Library

In order to fuzz a library, we employ the use of fuzzing drivers, or *harnesses*.

While some programs can simply be executed, libraries depend on other programs to invoke their functionality. Think of libraries as a toolbox, but someone still needs to actually pick up the tools and use them.

Fuzzing harnesses invoke functions of the library and pass fuzzer-generated inputs to the library in a meaningful way

Fuzzing Harness Example

```
#include <stdio.h>
#include <errno.h>
#include <unistd.h>
#include "ucl.h"

int LLVMFuzzerTestOneInput(const uint8_t *data, size_t size) {
    // If size is 0 we need a null-terminated string.
    // We dont null-terminate the string and by the design
    // of the API passing 0 as size with non null-terminated string
    // gives undefined behavior.
    if(size==0){
        return 0;
    }
    struct ucl_parser *parser;
    parser = ucl_parser_new(0);

    ucl_parser_add_string(parser, (char *)data, size);

    if (ucl_parser_get_error(parser) != NULL) {
        return 0;
    }

    ucl_parser_free (parser);
    return 0;
}
```

libucl

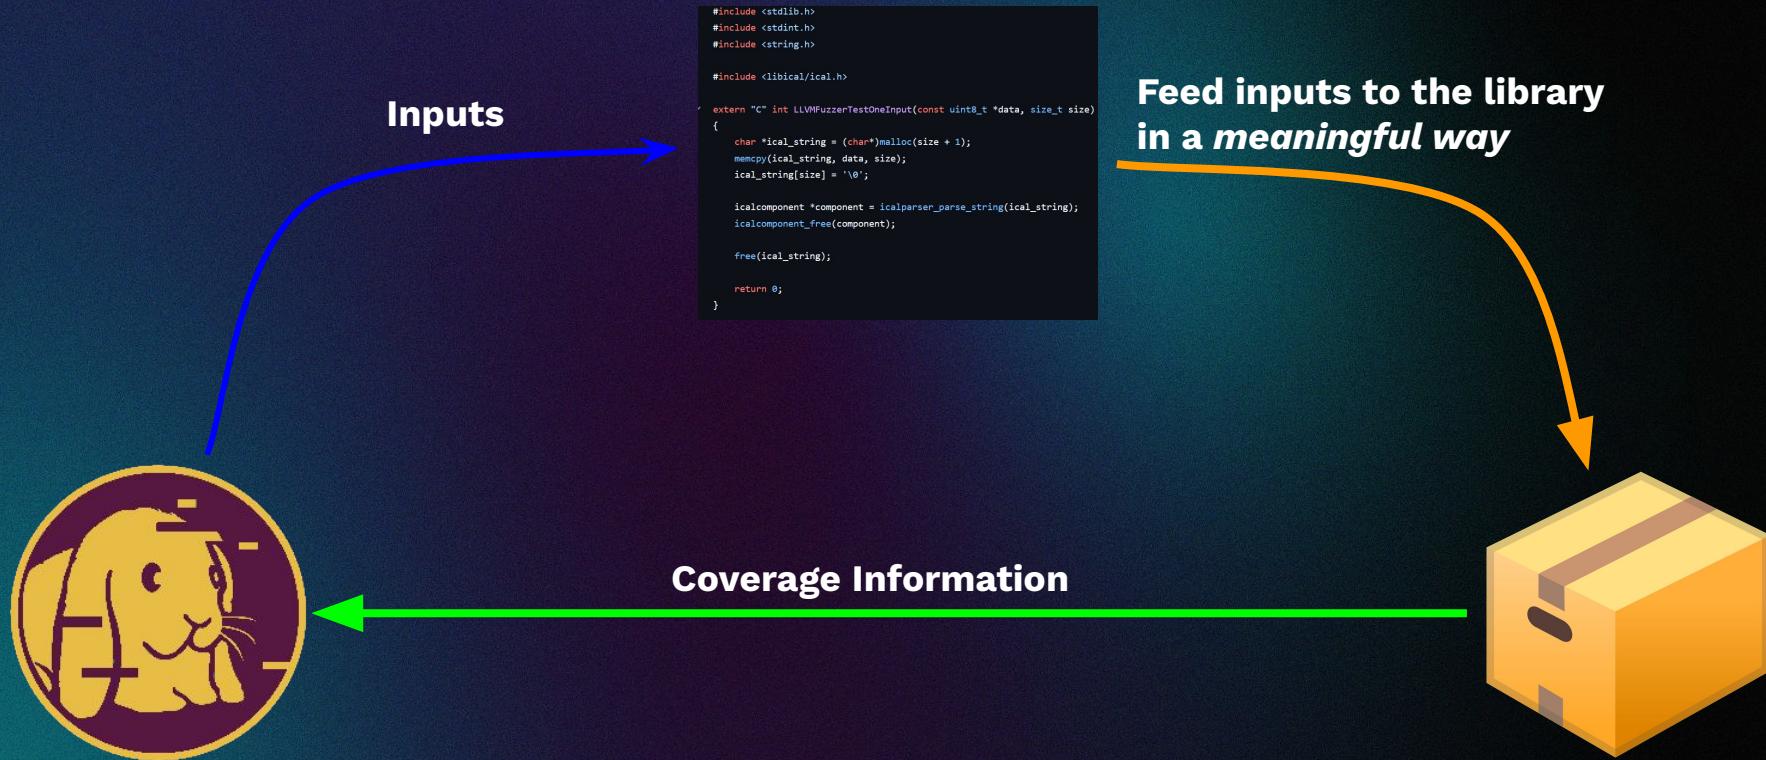
ucl_parser_new

ucl_parser_add_string

...



Library Fuzzing Overview



Let's Get Technical

These next slides describe how to build, instrument, and fuzz a library. Familiarity with build systems(CMake, autoconf, etc.), program compilation(clang), and the linux environment is recommended.

If not, don't worry! I've provided a demo of this build process here.

Google's OSS-Fuzz project is a great resource for library harnessing. Harnesses, library build scripts, and compilation scripts can all be found there.

<https://github.com/google/oss-fuzz/tree/master/projects>

Step 1: Build the Library

Goal: Build the library with AFLplusplus instrumentation

Depending on the build system, set the compiler to be AFL's own compiler:

- CMake: -DCMAKE_C_COMPILER=afl-clang-fast -DCMAKE_CXX_COMPILER=afl-clang-fast++
- Autoconf: CC=afl-clang-fast CXX=afl-clang-fast++

Step 2: Compilation

Goal: Compile the harness of interest and link it with the *instrumented* library

This varies between libraries and harnesses. Pretty much will need to use the -I, -L, and -l flags. For example, here's how you compile the harness in the c-ares demo.

```
afl-clang-fast -o harness.out harness.c  
-I lib_fuzz/include -I lib_fuzz/ -L lib_fuzz/.libs/ -lcares -fsanitize=address
```

Step 3: Start Fuzzing!

Goal: Run the fuzzing engine on your harness and fuzz the library!

Typical command to begin fuzzing:

```
afl-fuzz -i <seeds directory> -o <output directory> ./<binary name> @@
```

- new edges on section should increase.
If not, something's wrong.
- If *total crashes* is nonzero, you may have found a bug!

american fuzzy lop ++2.60d (libpng) [explore] {0}		overall results
process timing	run time : 0 days, 0 hrs, 4 min, 43 sec	cycles done : 0
	last new path : 0 days, 0 hrs, 0 min, 0 sec	total paths : 211
	last uniq crash : none seen yet	uniq crashes : 0
	last uniq hang : none seen yet	uniq hangs : 0
cycle progress	map coverage	
	now processing : 127.0 (60.2%)	map density : 0.24% / 1.67%
	paths timed out : 0 (0.00%)	count coverage : 1.81 bits/tuple
stage progress	findings in depth	
	now trying : arith 16/8	favored paths : 108 (51.18%)
	stage execs : 102/6963 (1.46%)	new edges on : 126 (59.72%)
	total execs : 1.39M	total crashes : 0 (0 unique)
	exec speed : 5147/sec	total tmouts : 0 (0 unique)
fuzzing strategy yields	path geometry	
	bit flips : 19/51.9k, 8/51.9k, 1/51.8k	levels : 3
	byte flips : 0/6488, 0/5262, 0/5210	pending : 160
	arithmetics : 7/295k, 1/114k, 0/52.1k	pend fav : 61
	known ints : 1/27.0k, 1/122k, 0/204k	own finds : 121
	dictionary : 0/0, 0/0, 3/257k	imported : 86
	havoc/rad : 23/114k, 0/0, 0/0	stability : 100.00%
	custom/rq : 0/0, 0/0, 45/9658, 12/15.2k	
	trim : 36.12%/3042, 18.23%	[cpu000: 31%]

Common Mistakes

new edges on section is not increasing:

- If your harness was not compiled with the *-static* flag, you need to set the LD_LIBRARY_PATH environment variable to point to the instrumented .so file.
 - Run `ldd ./<binary_name>`. This will give you the path of each library the harness is talking to. If the instrumented library is not in this path, set the LD_LIBRARY_PATH.
- Make sure you are linking your harness against the correct library.
- Your library is not actually instrumented with AFL.
 - To make sure the library is instrumented, find the .so file(the shared library). Run `objdump -D <path to .so file> | grep afl`. If nothing comes up the library is not instrumented!

Immediately getting a ton of crashes:

- You either found a bug really quick or your harness is incorrect.

afl-clang-fast not found:

- AFL isn't in your \$PATH. Either add it to the path or provide the absolute path to the binary

DEMO(~10 mins)

https://github.com/gabe-sherman/cybersec_fuzzing_demo

You need a Linux environment, or use WSL(Windows subsystem for linux)

- Fuzz a basic library and observe what a good fuzzing trial looks like
- Find a vulnerability in c-ares that allowed arbitrary code execution in < 5 seconds

CVE-2016-5180 Detail

MODIFIED

This CVE record has been updated after NVD enrichment efforts were completed. Enrichment data supplied by the NVD may require amendment due to these changes.

Current Description

Heap-based buffer overflow in the ares_create_query function in c-ares 1.x before 1.12.0 allows remote attackers to cause a denial of service (out-of-bounds write) or possibly execute arbitrary code via a hostname with an escaped trailing dot.

[View Analysis Description](#)

Metrics CVSS Version 4.0 CVSS Version 3.x CVSS Version 2.0

NVD enrichment efforts reference publicly available information to associate vector strings. CVSS information contributed by other sources is also displayed.

CVSS 3.x Severity and Vector Strings:

NVD NIST: NVD Base Score: 9.8 CRITICAL Vector: CVSS:3.1/AV:N/AC:L/PR:N/UI:N/S:U/C:H/I:H/A:H

Use the Internet!

There are a ton of blog posts out there describing the different aspects of fuzzing and harnessing. These are some of the best resources for learning.

<https://bushido-sec.com/index.php/2025/01/03/fuzzing-harness-guide/>

<https://google.github.io/clusterfuzz/setting-up-fuzzing/heartbleed-example/>

<https://www.mayhem.security/blog/firmware-fuzzing-101>

<https://www.fuzzingbook.org/>

<https://trustfoundry.net/2020/01/22/introduction-to-triaging-fuzzer-generated-crashes/>

Join Research!

Network Security



Cryptography



Mobile Security



Fuzzing



And more ... (AI security, systems security, exploits, cyber-physical systems, etc.)

Questions?